

y1bxijjmt

November 6, 2025

```
[ ]: import numpy as np
import pandas as pd
import seaborn as sns; sns.set(style="ticks", color_codes=True)
import matplotlib.pyplot as plt
import plotly.graph_objs as go
from plotly.subplots import make_subplots
import seaborn as sns
import plotly.express as px
%matplotlib inline

from plotly.offline import download_plotlyjs, init_notebook_mode, plot, iplot
init_notebook_mode(connected = True)
```

```
[12]: df = pd.read_csv(r'C://Users//anand//Downloads//DataCoSupplyChainDataset.csv',encoding='ISO-8859-1')

pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
pd.set_option('display.width', 1000)

df.head()
```

```
[12]:      Type Days for shipping (real) Days for shipment (scheduled) Benefit
per order Sales per customer Delivery Status Late_delivery_risk Category Id
Category Name Customer City Customer Country Customer Email Customer Fname
Customer Id Customer Lname Customer Password Customer Segment Customer State
Customer Street Customer Zipcode Department Id Department Name Latitude
Longitude Market Order City Order Country Order Customer Id order date
(DateOrders) Order Id Order Item Cardprod Id Order Item Discount Order Item
Discount Rate Order Item Id Order Item Product Price Order Item Profit Ratio
Order Item Quantity Sales Order Item Total Order Profit Per Order Order
Region Order State Order Status Order Zipcode Product Card Id
Product Category Id Product Description Product
Image Product Name Product Price Product Status shipping date (DateOrders)
Shipping Mode
0 DEBIT 3 4
91.250000 314.640015 Advance shipping 0 73
```

Sporting Goods	Caguas	Puerto Rico	XXXXXXXXXX	Cally
20755	Holloway	XXXXXXXXXX	Consumer	PR 5365
Noble Nectar Island		725.0	2	Fitness 18.251453
-66.037056	Pacific Asia	Bekasi	Indonesia	20755
1/31/2018 22:56	77202		1360	13.110000
0.04	180517		327.75	0.29
1 327.75	314.640015		91.250000	Southeast Asia Java
Occidental	COMPLETE	NaN	1360	73
NaN http://images.acmesports.sports/Smart+watch		Smart watch		327.75
0 2/3/2018 22:56	Standard Class			
1 TRANSFER		5		4
-249.089996	311.359985	Late delivery		1
73 Sporting Goods	Caguas	Puerto Rico	XXXXXXXXXX	Irene
19492	Luna	XXXXXXXXXX	Consumer	PR
2679 Rustic Loop		725.0	2	Fitness 18.279451
-66.037064	Pacific Asia	Bikaner	India	19492
1/13/2018 12:27	75939		1360	16.389999
0.05	179254		327.75	-0.80
1 327.75	311.359985		-249.089996	South Asia
Rajastán	PENDING	NaN	1360	73
NaN http://images.acmesports.sports/Smart+watch		Smart watch		327.75
0 1/18/2018 12:27	Standard Class			
2 CASH		4		4
-247.779999	309.720001	Shipping on time		0
73 Sporting Goods	San Jose	EE. UU.	XXXXXXXXXX	Gillian
19491	Maldonado	XXXXXXXXXX	Consumer	CA 8510
Round Bear Gate		95125.0	2	Fitness 37.292233
-121.881279	Pacific Asia	Bikaner	India	19491
1/13/2018 12:06	75938		1360	18.030001
0.06	179253		327.75	-0.80
1 327.75	309.720001		-247.779999	South Asia
Rajastán	CLOSED	NaN	1360	73
NaN http://images.acmesports.sports/Smart+watch		Smart watch		327.75
0 1/17/2018 12:06	Standard Class			
3 DEBIT		3		4
22.860001	304.809998	Advance shipping		0 73
Sporting Goods	Los Angeles	EE. UU.	XXXXXXXXXX	Tana
19490	Tate	XXXXXXXXXX	Home Office	CA
3200 Amber Bend		90027.0	2	Fitness 34.125946
-118.291016	Pacific Asia	Townsville	Australia	19490
1/13/2018 11:45	75937		1360	22.940001
0.07	179252		327.75	0.08
1 327.75	304.809998		22.860001	Oceania
Queensland	COMPLETE	NaN	1360	73
NaN http://images.acmesports.sports/Smart+watch		Smart watch		327.75
0 1/16/2018 11:45	Standard Class			
4 PAYMENT		2		4

134.210007	298.250000	Advance shipping	0	
73 Sporting Goods	Caguas	Puerto Rico	XXXXXXXXXX	Orli
19489 Hendricks	XXXXXXXXXX	Corporate	PR	8671
Iron Anchor Corners	725.0	2	Fitness	18.253769
-66.037048 Pacific Asia	Townsville	Australia	19489	
1/13/2018 11:24	75936	1360	29.500000	
0.09	179251	327.75	0.45	
1 327.75	298.250000	134.210007	Oceania	
Queensland PENDING_PAYMENT	NaN	1360	73	
Nan http://images.acmesports.sports/Smart+watch	Smart watch	327.75		
0 1/15/2018 11:24 Standard Class				

[15]: df.tail()

Type Days for shipping (real)	Days for shipment (scheduled)																			
Benefit per order	Sales per customer	Delivery Status	Late_delivery_risk																	
Category Id	Category Name	Customer City	Customer Country	Customer Email	Customer Fname	Customer Id	Customer Lname	Customer Password	Customer Segment	Customer State	Customer Street	Customer Zipcode	Department Id	Department Name	Latitude	Longitude	Market Order	City Order	Country Order	
Customer Id	order date (DateOrders)	Order Id	Order Item	Cardprod Id	Order Item Discount	Order Item Discount Rate	Order Item Id	Order Item Product Price	Order Item Profit Ratio	Order Item Quantity	Sales	Order Item Total	Order Profit Per Order	Order Region	Order State	Order Status	Order Zipcode	Product Card Id	Product Category Id	Product Description
Product Image	Product Name	Product Price	Product Status	shipping date (DateOrders) \																
180514 CASH	4	4																		
40.000000	399.980011	Shipping on time	0	45																
Fishing Brooklyn	EE. UU.	XXXXXXXXXX	Maria																	
1005 Peterson	XXXXXXXXXX	Home Office	NY																	
1322 Broad Glade	11207.0	7	Fan Shop	40.640930																
-73.942711 Pacific Asia	Shanghái	China	1005																	
1/16/2016 3:40	26043	1004	0.0																	
0.00 65177	399.980011	399.980011	0.10																	
1 399.980011	399.980011	40.000000	Eastern Asia																	
Shanghái CLOSED	NaN	1004	45																	
Nan http://images.acmesports.sports/Field+%	26+Stre...	Field & Stream Sportsman																		
16 Gun Fire Safe	399.980011	0	1/20/2016 3:40																	
180515 DEBIT	3	2																		
-613.770019	395.980011	Late delivery	1																	
45 Fishing Bakersfield	EE. UU.	XXXXXXXXXX	Ronald																	
9141 Clark	XXXXXXXXXX	Corporate	CA	7330																
Broad Apple Moor	93304.0	7	Fan Shop	35.362545																
-119.018700 Pacific Asia	Hirakata	Japón	9141																	
1/16/2016 1:34	26037	1004	4.0																	
0.01 65161	399.980011	399.980011	-1.55																	

1	399.980011	395.980011	-613.770019	Eastern Asia	
Osaka	COMPLETE	NaN	1004	45	
NaN	http://images.acmesports.sports/Field+%	Stre...	Field & Stream Sportsman		
16	Gun Fire Safe	399.980011	0	1/19/2016 1:34	
180516	TRANSFER		5	4	
141.110001		391.980011	Late delivery	1	
45	Fishing	Bristol	EE. UU.	XXXXXXX	John
291	Smith	XXXXXXXXXX	Corporate	CT	97
Burning Landing		6010.0	7	Fan Shop	41.629959
-72.967155	Pacific Asia	Adelaide	Australia		291
1/15/2016 21:00	26024		1004		8.0
0.02	65129		399.980011		0.36
1	399.980011	391.980011	141.110001	Oceania	Australia
del Sur	PENDING	NaN	1004		45
NaN	http://images.acmesports.sports/Field+%	Stre...	Field & Stream Sportsman		
16	Gun Fire Safe	399.980011	0	1/20/2016 21:00	
180517	PAYMENT		3		4
186.229996		387.980011	Advance shipping		0
45	Fishing	Caguas	Puerto Rico	XXXXXXX	Mary
2813	Smith	XXXXXXXXXX	Consumer	PR	2585
Silent Autumn Landing		725.0	7	Fan Shop	
18.213350	-66.370575	Pacific Asia	Adelaide	Australia	2813
1/15/2016 20:18	26022		1004		12.0
0.03	65126		399.980011		0.48
1	399.980011	387.980011	186.229996	Oceania	Australia
del Sur	PENDING_PAYMENT	NaN	1004		45
NaN	http://images.acmesports.sports/Field+%	Stre...	Field & Stream Sportsman		
16	Gun Fire Safe	399.980011	0	1/18/2016 20:18	
180518	PAYMENT		4		4
168.949997		383.980011	Shipping on time		0
45	Fishing	Caguas	Puerto Rico	XXXXXXX	Andrea
7547	Ortega	XXXXXXXXXX	Consumer	PR	
697 Little Meadow		725.0	7	Fan Shop	18.290380
-66.370613	Pacific Asia	Nagercoil	India		7547
1/15/2016 18:54	26018		1004		16.0
0.04	65113		399.980011		0.44
1	399.980011	383.980011	168.949997	South Asia	
Tamil Nadu	PENDING_PAYMENT	NaN	1004		45
NaN	http://images.acmesports.sports/Field+%	Stre...	Field & Stream Sportsman		
16	Gun Fire Safe	399.980011	0	1/19/2016 18:54	

Shipping Mode

180514	Standard Class
180515	Second Class
180516	Standard Class
180517	Standard Class
180518	Standard Class

```
[13]: df.shape
```

```
[13]: (180519, 53)
```

```
[14]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 180519 entries, 0 to 180518
Data columns (total 53 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Type             180519 non-null   object  
 1   Days for shipping (real)    180519 non-null   int64  
 2   Days for shipment (scheduled) 180519 non-null   int64  
 3   Benefit per order        180519 non-null   float64 
 4   Sales per customer       180519 non-null   float64 
 5   Delivery Status         180519 non-null   object  
 6   Late_delivery_risk      180519 non-null   int64  
 7   Category Id            180519 non-null   int64  
 8   Category Name          180519 non-null   object  
 9   Customer City          180519 non-null   object  
 10  Customer Country        180519 non-null   object  
 11  Customer Email         180519 non-null   object  
 12  Customer Fname         180519 non-null   object  
 13  Customer Id            180519 non-null   int64  
 14  Customer Lname         180511 non-null   object  
 15  Customer Password       180519 non-null   object  
 16  Customer Segment        180519 non-null   object  
 17  Customer State          180519 non-null   object  
 18  Customer Street         180519 non-null   object  
 19  Customer Zipcode        180516 non-null   float64 
 20  Department Id          180519 non-null   int64  
 21  Department Name         180519 non-null   object  
 22  Latitude                180519 non-null   float64 
 23  Longitude               180519 non-null   float64 
 24  Market                  180519 non-null   object  
 25  Order City              180519 non-null   object  
 26  Order Country            180519 non-null   object  
 27  Order Customer Id        180519 non-null   int64  
 28  order date (DateOrders) 180519 non-null   object  
 29  Order Id                 180519 non-null   int64  
 30  Order Item Cardprod Id   180519 non-null   int64  
 31  Order Item Discount       180519 non-null   float64 
 32  Order Item Discount Rate 180519 non-null   float64 
 33  Order Item Id            180519 non-null   int64  
 34  Order Item Product Price 180519 non-null   float64 
 35  Order Item Profit Ratio   180519 non-null   float64
```

```

36 Order Item Quantity           180519 non-null   int64
37 Sales                         180519 non-null   float64
38 Order Item Total              180519 non-null   float64
39 Order Profit Per Order       180519 non-null   float64
40 Order Region                  180519 non-null   object
41 Order State                   180519 non-null   object
42 Order Status                  180519 non-null   object
43 Order Zipcode                 24840 non-null    float64
44 Product Card Id              180519 non-null   int64
45 Product Category Id          180519 non-null   int64
46 Product Description           0 non-null        float64
47 Product Image                 180519 non-null   object
48 Product Name                  180519 non-null   object
49 Product Price                 180519 non-null   float64
50 Product Status                180519 non-null   int64
51 shipping date (DateOrders)   180519 non-null   object
52 Shipping Mode                 180519 non-null   object
dtypes: float64(15), int64(14), object(24)
memory usage: 73.0+ MB

```

[4]: df.columns

```

[4]: Index(['Type', 'Days for shipping (real)', 'Days for shipment (scheduled)',
'Benefit per order', 'Sales per customer', 'Delivery Status',
'Late_delivery_risk', 'Category Id', 'Category Name', 'Customer City', 'Customer
Country', 'Customer Email', 'Customer Fname', 'Customer Id', 'Customer Lname',
'Customer Password', 'Customer Segment', 'Customer State', 'Customer Street',
'Customer Zipcode', 'Department Id', 'Department Name', 'Latitude', 'Longitude',
'Market', 'Order City', 'Order Country', 'Order Customer Id', 'order date
(DateOrders)', 'Order Id', 'Order Item Cardprod Id', 'Order Item Discount',
'Order Item Discount Rate', 'Order Item Id', 'Order Item Product Price', 'Order
Item Profit Ratio', 'Order Item Quantity', 'Sales', 'Order Item Total', 'Order
Profit Per Order', 'Order Region', 'Order State', 'Order Status', 'Order
Zipcode', 'Product Card Id', 'Product Category Id', 'Product Description',
'Product Image', 'Product Name', 'Product Price', 'Product Status',
'shipping date (DateOrders)', 'Shipping Mode'],
dtype='object')

```

deleting duplicate and not reated features

```

[5]: data=df.copy()
FeatureList=['Type', 'Benefit per order', 'Sales per customer',
             'Delivery Status', 'Late_delivery_risk', 'Category Name', 'Customer
             ↴City', 'Customer Country',
             'Customer Id', 'Customer Segment',
             'Customer State', 'Customer Zipcode', 'Department Name', 'Latitude', ↴
             ↴'Longitude',

```

```

        'Market', 'Order City', 'Order Country', 'Order Customer Id', 'order_
        ↵date (DateOrders)', 'Order Id',
        'Order Item Cardprod Id', 'Order Item Discount', 'Order Item Discount_
        ↵Rate', 'Order Item Id',
        'Order Item Product Price', 'Order Item Profit Ratio', 'Order Item_
        ↵Quantity', 'Sales', 'Order Item Total',
        'Order Profit Per Order', 'Order Region', 'Order State', 'Order_
        ↵Status', 'Order Zipcode', 'Product Card Id',
        'Product Category Id', 'Product Description', 'Product Image',_
        ↵'Product Name', 'Product Price', 'Product Status',
        'shipping date (DateOrders)', 'Shipping Mode']
    
```

df1=df[FeatureList]
df1.head()

[5]:

	Type	Benefit per order	Sales per customer	Delivery Status
Late_delivery_risk	Category Name	Customer City	Customer Country	Customer Id
Customer Segment	Customer State	Customer Zipcode	Department Name	Latitude
Longitude	Market	Order City	Order Country	Order Customer Id
order date (DateOrders)	Order Id	Order Item Cardprod Id	Order Item Discount	Order Item Discount Rate
Order Item Id	Order Item Product Price	Order Item Profit Ratio	Order Item Total	Order Profit Per Order
Order Item Quantity	Sales	Order Profit Per Order	Order Region	Order Status
Order Region	Order Status	Order Zipcode	Product Card Id	Product Category Id
Product Description	Product Image	Product Name	Product Price	Product Status
Shipping Mode	shipping date (DateOrders)			
0 DEBIT	91.250000	314.640015	Advance shipping	
0 Sporting Goods	Caguas	Puerto Rico	20755	Consumer
PR 725.0	Fitness	18.251453	-66.037056	Pacific Asia
Bekasi Indonesia		20755	1/31/2018 22:56	77202
1360 13.110000		0.04	180517	
327.75 0.29		1 327.75	314.640015	
91.250000 Southeast Asia	Java Occidental		COMPLETE	NaN
1360 73		Nan		
http://images.acmesports.sports/Smart+watch	Smart watch		327.75	
0 2/3/2018 22:56 Standard Class				
1 TRANSFER -249.089996	311.359985	Late delivery		
1 Sporting Goods	Caguas	Puerto Rico	19492	Consumer
PR 725.0	Fitness	18.279451	-66.037064	Pacific Asia
Bikaner India		19492	1/13/2018 12:27	75939
1360 16.389999		0.05	179254	
327.75 -0.80		1 327.75	311.359985	
-249.089996 South Asia	Rajast n		PENDING	NaN
1360 73		Nan		
http://images.acmesports.sports/Smart+watch	Smart watch		327.75	
0 1/18/2018 12:27 Standard Class				
2 CASH -247.779999	309.720001	Shipping on time		

```

0 Sporting Goods San Jose EE. UU. 19491 Consumer
CA 95125.0 Fitness 37.292233 -121.881279 Pacific Asia
Bikaner India 19491 1/13/2018 12:06 75938
1360 18.030001 0.06 179253
327.75 -0.80 1 327.75 309.720001
-247.779999 South Asia Rajast n CLOSED NaN
1360 73 NaN
http://images.acmesports.sports/Smart+watch Smart watch 327.75
0 1/17/2018 12:06 Standard Class
3 DEBIT 22.860001 304.809998 Advance shipping
0 Sporting Goods Los Angeles EE. UU. 19490 Home Office
CA 90027.0 Fitness 34.125946 -118.291016 Pacific Asia
Townsville Australia 19490 1/13/2018 11:45 75937
1360 22.940001 0.07 179252
327.75 0.08 1 327.75 304.809998
22.860001 Oceania Queensland COMPLETE NaN
1360 73 NaN
http://images.acmesports.sports/Smart+watch Smart watch 327.75
0 1/16/2018 11:45 Standard Class
4 PAYMENT 134.210007 298.250000 Advance shipping
0 Sporting Goods Caguas Puerto Rico 19489 Corporate
PR 725.0 Fitness 18.253769 -66.037048 Pacific Asia
Townsville Australia 19489 1/13/2018 11:24 75936
1360 29.500000 0.09 179251
327.75 0.45 1 327.75 298.250000
134.210007 Oceania Queensland PENDING_PAYMENT NaN
1360 73 NaN
http://images.acmesports.sports/Smart+watch Smart watch 327.75
0 1/15/2018 11:24 Standard Class

```

1 Data Visualization

2 delivery status

```
[6]: data_delivery_status=df1.groupby(['Delivery Status'])['Order Id'].count().
    ↪reset_index(name='Number of Orders').sort_values(by= 'Number of Orders', ↪
    ↪ascending= False)
px.bar(x=data_delivery_status['Delivery Status'] , ↪
    ↪y=data_delivery_status['Number of Orders'] , ↪
    ↪color=data_delivery_status['Number of Orders'],
        labels = { 'Delivery Status': 'Delivery Status', 'Number of Orders': ↪
    ↪'Number of Orders'})
```

```
[7]: data_delivery_status_region=df1.groupby(['Delivery Status', 'Order ↪
    ↪Region'])['Order Id'].count().reset_index(name='Number of Orders').
    ↪sort_values(by= 'Number of Orders', ascending= False)
```

```
px.bar(data_delivery_status_region, x='Delivery Status', y='Number of Orders' ▾
    ↵, color='Order Region',
    )
```

Central America and Western Europa have the most lated delivered Orders

3 Top 20 Customers regarding the quantity of orders

```
[8]: df1['Customer_ID_STR']=df1['Customer Id'].astype(str)

data_customers=df1.groupby(['Customer_ID_STR'])['Order Id'].count().
    ↵reset_index(name='Number of Orders').sort_values(by= 'Number of Orders', ▾
    ↵ascending= False)
px.bar(data_customers.head(20),x='Number of Orders', y='Customer_ID_STR' , ▾
    ↵color='Number of Orders'      )
```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:1:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

4 Top 20 Customers regarding profit of all orders

```
[9]: df1['Customer_ID_STR']=df1['Customer Id'].astype(str)

data_customers_profit=df1.groupby(['Customer_ID_STR'])['Order Profit Per Order'].sum().reset_index(name='Profit of Orders').sort_values(by= 'Profit of Orders', ascending= False)
px.bar(data_customers_profit.head(20),x='Profit of Orders', y='Customer_ID_STR' , ▾
    ↵, color='Profit of Orders'      )
```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:1:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

5 Customer Segment

```
[10]: #Customer Segments
data_Customer_Segment=df1.groupby(['Customer Segment'])['Order Id'].count().
    ↪reset_index(name='Number of Orders').sort_values(by= 'Number of Orders',�
    ↪ascending= False)
px.pie(data_Customer_Segment, values='Number of Orders', names= 'Customer�
    ↪Segment' , title= 'Number of Orders of different Customer Segments',
        width=600 , height=600 , color_discrete_sequence = px.colors.sequential.
    ↪RdBu)
```

6 Category

```
[11]: #Category Name
data_Category_Name=df1.groupby(['Category Name'])['Order Id'].count().
    ↪reset_index(name='Number of Orders').sort_values(by= 'Number of Orders',�
    ↪ascending= True)
px.bar(data_Category_Name, x='Number of Orders',y = 'Category Name',color=
    ↪='Number of Orders')
```

7 Geo Features

```
[12]: data_Region=df1.groupby(['Order Region'])['Order Id'].count().
    ↪reset_index(name='Number of Orders').sort_values(by= 'Number of Orders',�
    ↪ascending= True)
px.bar(data_Region, x='Number of Orders',y = 'Order Region',color ='Number of�
    ↪Orders')
```

```
[13]: data_countries=df1.groupby(['Order Country'])['Order Id'].count().
    ↪reset_index(name='Number of Orders').sort_values(by= 'Number of Orders',�
    ↪ascending= True)
px.bar(data_countries.head(20), x='Number of Orders',y = 'Order Country',color=
    ↪='Number of Orders')
```

```
[14]: df_geo=df1.groupby(['Order Country', 'Order City'])['Order Profit Per Order'].
    ↪sum().reset_index(name='Profit of Orders').sort_values(by= 'Profit of Orders', ascending= False)
df_geo
```

```
[14]:          Order Country      Order City  Profit of Orders
3260  Rep blica Dominicana  Santo Domingo      51111.670019
1492        Estados Unidos  New York City      47889.759868
2152            Honduras    Tegucigalpa      40973.640056
1430        Estados Unidos  Los Angeles      38014.360024
2837        Nicaragua       Managua      34319.950107
...
738            China         Dalian      -1588.609972
3        Afganist n       Kandahar      -1681.830001
3110        Reino Unido     Dudley      -1742.079996
2463           Italia       Cerignola      -2212.530012
125        Alemania       Pulheim      -3152.150012
```

[3665 rows x 3 columns]

```
[15]: fig = px.choropleth(df_geo ,  locationmode='country names', locations='Order Country',
                        color='Profit of Orders', # lifeExp is a column of data
                        hover_name='Order Country',
                        #hover_data ='Order City',
                        color_continuous_scale=px.colors.sequential.Plasma)

fig.show()
```

8 Sales Analysis

```
[16]: #Order Country
df_sales_country=df1.groupby(['Order Country'])['Sales'].sum().
    reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders', ascending= False)
px.bar(df_sales_country.head(10), x='Sales of Orders',y = 'Order Country',color= 'Sales of Orders')
```

[]:

```
[17]: #Order Country
df_sales_country=df1.groupby(['Order Country'])['Sales'].sum().
    reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders', ascending= False)
px.bar(df_sales_country.head(10), x='Sales of Orders',y = 'Order Country',color= 'Sales of Orders')
```

```
[18]: #Product
df_sales_country=df1.groupby(['Product Name'])['Sales'].sum().
    reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders', ascending= False)
```

```
px.bar(df_sales_country.head(10), x='Sales of Orders',y = 'Product Name',color= 'Sales of Orders')
```

```
[19]: #Product and delivery status  
df_sales_pd=df1.groupby(['Product Name', 'Delivery Status'])['Sales'].sum().  
    ↪reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders',  
    ↪ascending= False)  
px.bar(df_sales_pd.head(10), x='Sales of Orders',y = 'Product Name',color= 'Delivery Status')
```

```
[20]: #Product and order region  
df_sales_pr=df1.groupby(['Product Name', 'Order Region'])['Sales'].sum().  
    ↪reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders',  
    ↪ascending= False)  
px.bar(df_sales_pr.head(10), x='Sales of Orders',y = 'Product Name',color= 'Order Region')
```

```
[21]: #'Category Name'  
df_sales_pr=df1.groupby(['Category Name'])['Sales'].sum().  
    ↪reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders',  
    ↪ascending= False)  
px.bar(df_sales_pr.head(10), x='Sales of Orders',y = 'Category Name',color= 'Sales of Orders')
```

```
[22]: #'Type of payment  
df_sales_pr=df1.groupby(['Type'])['Sales'].sum().reset_index(name='Sales of  
    ↪Orders').sort_values(by= 'Sales of Orders', ascending= False)  
px.bar(df_sales_pr.head(10), x='Sales of Orders',y = 'Type',color = 'Sales of  
    ↪Orders')
```

```
[23]: df_sales_tp=df1.groupby(['Type', 'Product Name'])['Sales'].sum().  
    ↪reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders',  
    ↪ascending= False)  
px.bar(df_sales_tp.head(10), x='Sales of Orders',y = 'Type',color = 'Product  
    ↪Name')
```

9 Date and sales analysis

```
[24]: import datetime as dt  
  
data_orderdate=df[['order date (DateOrders)', 'Sales']]  
data_orderdate['order_date'] = pd.to_datetime(data_orderdate['order date  
    ↪(DateOrders)'])
```

```
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:4:  
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
[25]: data_orderdate["Quarter"] = data_orderdate['order_date'].dt.quarter  
data_orderdate["Month"] = data_orderdate['order_date'].dt.month  
data_orderdate["year"] = data_orderdate['order_date'].dt.year
```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:1:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:2:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:3:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
[26]: data_orderdate['YearStr']=data_orderdate['year'].astype(str)
```

```

df_sales_year=data_orderdate.groupby(['YearStr'])['Sales'].sum().
    ↪reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders',u
    ↪ascending= False)
px.bar(df_sales_year, x='Sales of Orders',y = 'YearStr',color ='Sales of u
    ↪Orders')

```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:1:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

[27]:

```

data_orderdate['QuarterStr']=data_orderdate['Quarter'].astype(str)
df_sales_quarter=data_orderdate.groupby(['YearStr','QuarterStr'])['Sales'].
    ↪sum().reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders',u
    ↪ascending= False)
px.bar(df_sales_quarter, x='Sales of Orders',y = 'QuarterStr',color ='YearStr')

```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:1:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

[28]:

```

data_orderdate['MonthStr']=data_orderdate['Month'].astype(str)
df_sales_m=data_orderdate.groupby(['QuarterStr', 'MonthStr'])['Sales'].sum().
    ↪reset_index(name='Sales of Orders').sort_values(by= 'Sales of Orders',u
    ↪ascending= False)
px.bar(df_sales_m, x='Sales of Orders',y = 'QuarterStr',color ='MonthStr')

```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:1:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

[]:

10 Forecasting

Predicting if an order is fraud or not

```
[29]: data=df1.copy()
data['SUSPECTED_FRAUD'] = np.where(data['Order Status'] == 'SUSPECTED_FRAUD', 1, 0)
```

```
[30]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
def Labelencoder_feature(x):
    le=LabelEncoder()
    x=le.fit_transform(x)
    return x
```

```
[31]: features=data.drop(columns=['SUSPECTED_FRAUD','Order Status'])
target=data['SUSPECTED_FRAUD']
```

```
[32]: features.isnull().sum()
```

```
[32]: Type                      0
Benefit per order              0
Sales per customer              0
Delivery Status                0
Late_delivery_risk              0
Category Name                  0
Customer City                  0
Customer Country                0
Customer Id                     0
Customer Segment                0
Customer State                  0
Customer Zipcode                3
Department Name                 0
Latitude                         0
Longitude                        0
Market                            0
Order City                       0
Order Country                     0
Order Customer Id                0
order date (DateOrders)          0
Order Id                          0
```

```

Order Item Cardprod Id          0
Order Item Discount            0
Order Item Discount Rate       0
Order Item Id                  0
Order Item Product Price       0
Order Item Profit Ratio        0
Order Item Quantity            0
Sales                          0
Order Item Total               0
Order Profit Per Order         0
Order Region                   0
Order State                     0
Order Zipcode                  155679
Product Card Id                0
Product Category Id            0
Product Description             180519
Product Image                   0
Product Name                    0
Product Price                   0
Product Status                  0
shipping date (DateOrders)     0
Shipping Mode                   0
Customer_ID_STR                0
dtype: int64

```

we can consider NaN values as a separate class using LabelEncoder

```
[33]: features=features.apply(Labelencoder_feature)
features.head()
```

```

[33]: Type Benefit per order Sales per customer Delivery Status
Late_delivery_risk Category Name Customer City Customer Country Customer Id
Customer Segment Customer State Customer Zipcode Department Name Latitude
Longitude Market Order City Order Country Order Customer Id order date
(DateOrders) Order Id Order Item Cardprod Id Order Item Discount Order Item
Discount Rate Order Item Id Order Item Product Price Order Item Profit Ratio
Order Item Quantity Sales Order Item Total Order Profit Per Order Order
Region Order State Order Zipcode Product Card Id Product Category Id
Product Description Product Image Product Name Product Price Product Status
shipping date (DateOrders) Shipping Mode Customer_ID_STR
0      1           18934           2568           0
0      40          66              1           20649           0
36      7           4           3624          4420           3           331
70      20649           5961          65749           114
426           4           180516           62
140           0           166           2568           18934
15      475          12733          114           47

```

0		78	78	62	0		
27149		3	11912				
1	3	2272		2559	1		
1		40	66		19386		0
36		7	4	5522	4419	3	391
69		19386		1147	64486		114
498			5	179253			62
44		0	166	2559			2272
13	841		100315	114			47
120339		78	78	62		0	
2209		3	10509				
2	0	2293		2555		3	
0		40	452		19385		0
5		941		9146	232	3	391
69		19385		1146	64485		114
529			6	179252			62
44		0	166	2555			2293
13	841		100316	114			47
120340		78	78	62		0	
1980		3	10508				
3	1	13638		2546		0	
0		40	285		19384		2
5		720		8467	663	3	3226
8		19384		1145	64484		114
600			7	179251			62
119		0	166	2546			13638
11	835		100317	114			47
120341		78	78	62		0	
1752		3	10507				
4	2	20599		2526		0	
0		40	66		19383		1
36		7		3783	4421	3	3226
8		19383		1144	64483		114
682			8	179250			62
156		0	166	2526			20599
11	835		100318	114			47
120342		78	78	62		0	
1528		3	10505				

[34]: #deleting features which high-correlated with other features to avoid
 ↵multicollinearity
`data1=pd.concat([features,target],axis=1)`

[35]: #deleting features which high-correlated with other features to avoid
 ↵multicollinearity
`corr = data1.corr()`

```

columns = np.full((corr.shape[0],), True, dtype=bool)
for i in range(corr.shape[0]):
    for j in range(i+1, corr.shape[0]):
        if corr.iloc[i,j] >= 0.8:
            if columns[j]:
                columns[j] = False
selected_columns = data1.columns[columns]
selected_columns

```

[35]: Index(['Type', 'Benefit per order', 'Sales per customer', 'Delivery Status', 'Late_delivery_risk', 'Category Name', 'Customer City', 'Customer Country', 'Customer Id', 'Customer Segment', 'Customer State', 'Customer Zipcode', 'Department Name', 'Latitude', 'Longitude', 'Market', 'Order City', 'Order Country', 'order date (DateOrders)', 'Order Id', 'Order Item Cardprod Id', 'Order Item Discount', 'Order Item Discount Rate', 'Order Item Product Price', 'Order Item Quantity', 'Order Region', 'Order State', 'Order Zipcode', 'Product Description', 'Product Image', 'Product Status', 'Shipping Mode', 'Customer_ID_STR', 'SUSPECTED_FRAUD'], dtype='object')

[36]: features1=features[['Type', 'Benefit per order', 'Sales per customer', 'Delivery Status', 'Late_delivery_risk', 'Category Name', 'Customer City', 'Customer Country', 'Customer Id', 'Customer Segment', 'Customer State', 'Customer Zipcode', 'Department Name', 'Latitude', 'Longitude', 'Market', 'Order City', 'Order Country', 'order date (DateOrders)', 'Order Id', 'Order Item Cardprod Id', 'Order Item Discount', 'Order Item Discount Rate', 'Order Item Product Price', 'Order Item Quantity', 'Order Region', 'Order State', 'Order Zipcode', 'Product Description', 'Product Image', 'Product Status', 'Shipping Mode', 'Customer_ID_STR']]

[37]: from scipy.stats import pearsonr
corre=pd.DataFrame()
for i in features1.columns:
 corre[i]= pearsonr(target, features1[i])

corre

/opt/conda/lib/python3.7/site-packages/scipy/stats/stats.py:3913:
PearsonRConstantInputWarning:

An input array is constant; the correlation coefficient is not defined.

```
[37]:      Type Benefit per order Sales per customer Delivery Status
Late_delivery_risk Category Name Customer City Customer Country Customer Id
Customer Segment Customer State Customer Zipcode Department Name Latitude
Longitude Market Order City Order Country order date (DateOrders) Order
Id Order Item Cardprod Id Order Item Discount Order Item Discount Rate Order
Item Product Price Order Item Quantity Order Region Order State Order
Zipcode Product Description Product Image Product Status Shipping Mode
Customer_ID_STR
0 0.202094 -0.002586 -0.000807 0.128768
-0.167158 -0.002104 -0.003561 0.006756 0.009375
-0.005935 0.005670 -0.005600 -0.000837 -0.004709
0.005100 -0.000293 -0.005889 -0.004928 0.003002 0.002222
-0.001109 0.002893 0.002271
0.000315 -0.000757 0.009077 -0.006105 -4.274182e-02
-5.615812e-02 -0.000741 NaN -0.004641 0.000490
1 0.000000 0.271923 0.731807 0.000000
0.000000 0.371352 0.130283 0.004098 0.000068
0.011686 0.016001 0.017337 0.722187 0.045427 0.030231
0.900798 0.012345 0.036290 0.202217 0.345187
0.637360 0.219059 0.334515
0.893383 0.747677 0.000115 0.009491 9.230621e-74
5.077551e-126 0.752963 NaN 0.048639 0.835012
```

```
[38]: corre1=corre.T
```

```
[39]: coore2= corre1.iloc[:,0].sort_values(ascending=False)
```

```
coore2
```

Type	0.202094
Delivery Status	0.128768
Customer Id	0.009375
Order Region	0.009077
Customer Country	0.006756
Customer State	0.005670
Longitude	0.005100
order date (DateOrders)	0.003002
Order Item Discount	0.002893
Order Item Discount Rate	0.002271
Order Id	0.002222
Customer_ID_STR	0.000490
Order Item Product Price	0.000315
Market	-0.000293
Product Image	-0.000741
Order Item Quantity	-0.000757

```

Sales per customer      -0.000807
Department Name         -0.000837
Order Item Cardprod Id -0.001109
Category Name           -0.002104
Benefit per order       -0.002586
Customer City           -0.003561
Shipping Mode            -0.004641
Latitude                 -0.004709
Order Country            -0.004928
Customer Zipcode         -0.005600
Order City                -0.005889
Customer Segment          -0.005935
Order State               -0.006105
Order Zipcode              -0.042742
Product Description        -0.056158
Late_delivery_risk        -0.167158
Product Status             NaN
Name: 0, dtype: float64

```

[40]: coore2.index

[40]: Index(['Type', 'Delivery Status', 'Customer Id', 'Order Region', 'Customer Country', 'Customer State', 'Longitude', 'order date (DateOrders)', 'Order Item Discount', 'Order Item Discount Rate', 'Order Id', 'Customer_ID_STR', 'Order Item Product Price', 'Market', 'Product Image', 'Order Item Quantity', 'Sales per customer', 'Department Name', 'Order Item Cardprod Id', 'Category Name', 'Benefit per order', 'Customer City', 'Shipping Mode', 'Latitude', 'Order Country', 'Customer Zipcode', 'Order City', 'Customer Segment', 'Order State', 'Order Zipcode', 'Product Description', 'Late_delivery_risk', 'Product Status'],
dtype='object')

omit all features which have less than +-0.004 correlation with target

[41]: new_features= ['Type', 'Delivery Status', 'Order Region', 'Customer Country',
 ← 'Customer State', 'Order Zipcode',
 ← 'Shipping Mode', 'Order Country', 'Customer Zipcode', 'Order
 ← City', 'Customer Segment', 'Order State',
 ← 'Late_delivery_risk', 'Product Description', 'Product Status']

[42]: #Feature Selection

```

# Feature Selection based on importance
from sklearn.feature_selection import f_regression
F_values, p_values = f_regression(features, target)

```

```
/opt/conda/lib/python3.7/site-
packages/sklearn/feature_selection/_univariate_selection.py:302: RuntimeWarning:
```

```
invalid value encountered in true_divide
```

```
[43]: import itertools
f_reg_results = [(i, v, z) for i, v, z in itertools.zip_longest(features.
    ↪columns, F_values, [%.3f' % p for p in p_values])]
f_reg_results=pd.DataFrame(f_reg_results, columns=['Variable','F_Value', ↪
    ↪'P_Value'])
```

```
[44]: f_reg_results=pd.DataFrame(f_reg_results, columns=['Variable','F_Value', ↪
    ↪'P_Value'])
f_reg_results = f_reg_results.sort_values(by=['P_Value'])
f_reg_results.P_Value= f_reg_results.P_Value.astype(float)
f_reg_results=f_reg_results[f_reg_results.P_Value<0.06]
f_reg_results
```

```
[44]:
```

	Variable	F_Value	P_Value
0	Type	7686.615869	0.000
18	Order Customer Id	15.868141	0.000
3	Delivery Status	3043.655705	0.000
4	Late_delivery_risk	5188.953532	0.000
33	Order Zipcode	330.383488	0.000
8	Customer Id	15.868141	0.000
31	Order Region	14.875551	0.000
36	Product Description	571.103804	0.000
7	Customer Country	8.239934	0.004
32	Order State	6.728235	0.009
16	Order City	6.260807	0.012
9	Customer Segment	6.358072	0.012
10	Customer State	5.802926	0.016
11	Customer Zipcode	5.662017	0.017
14	Longitude	4.696175	0.030
17	Order Country	4.383497	0.036
13	Latitude	4.002769	0.045
42	Shipping Mode	3.887812	0.049
41	shipping date (DateOrders)	3.875811	0.049

```
[45]: f_reg_list=f_reg_results.Variable.values
f_reg_list
```

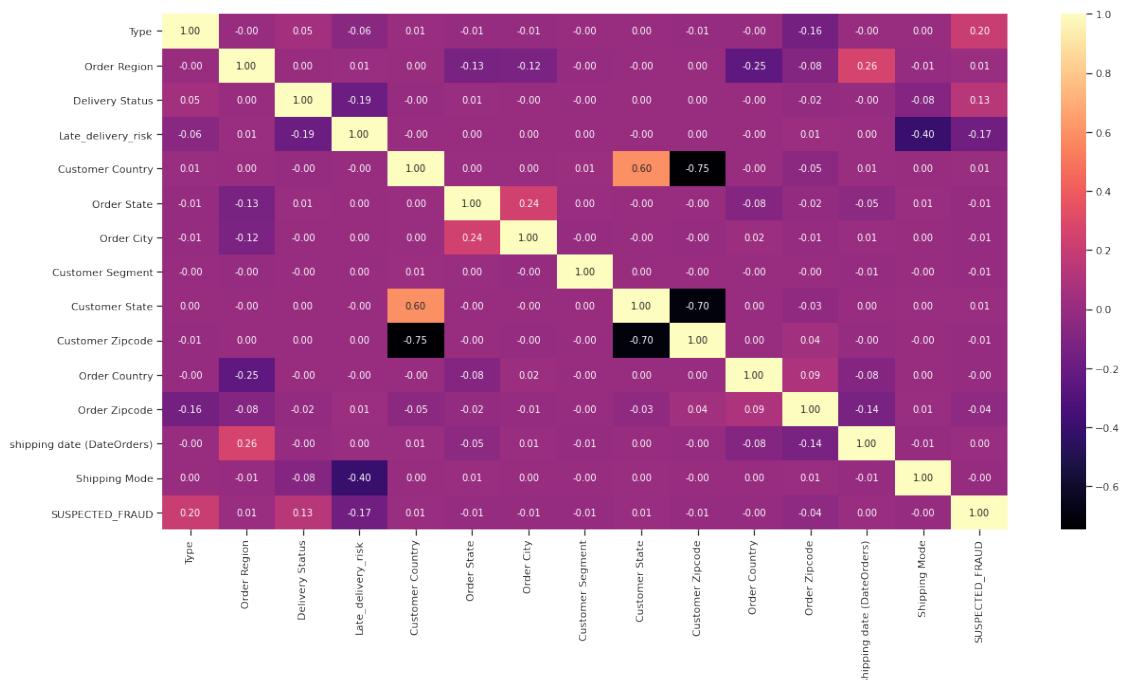
```
[45]: array(['Type', 'Order Customer Id', 'Delivery Status',
       'Late_delivery_risk', 'Order Zipcode', 'Customer Id',
       'Order Region', 'Product Description', 'Customer Country',
       'Order State', 'Order City', 'Customer Segment', 'Customer State',
       'Customer Zipcode', 'Longitude', 'Order Country', 'Latitude',
       'Shipping Mode', 'shipping date (DateOrders)'], dtype=object)
```

```
[46]: #final features list is both f_ref_list and new_feature
final_features=features[['Type', 'Order Region', 'Delivery Status', ↴
    'Late_delivery_risk',
    'Customer Country', 'Order State', 'Order City',
    'Customer Segment', 'Customer State', 'Customer Zipcode',
    'Order Country', 'Order Zipcode', 'shipping date (DateOrders)',
    'Shipping Mode']]
```

```
[47]: final_data=pd.concat([final_features, target], axis=1)
```

```
[48]: fig = plt.figure(figsize=(20,10))
sns.heatmap(final_data.corr(), annot = True, fmt = '.2f', cmap = 'magma')
```

[48]: <AxesSubplot:>



customer ZipCode and Customer state have high correlation with Customer Country. we can omit them and keep only customer country

```
[49]: final_features2=final_features.drop(columns=['Customer State', 'Customer Zipcode'])
```

```
[50]: from sklearn.pipeline import make_pipeline
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score, cross_validate
```

```

from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier
from sklearn.metrics import *
from sklearn import metrics
from sklearn.metrics import classification_report

```

[51]:

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(final_features2, target,
    test_size = 0.2, random_state = 42)

```

I use several classification models using also CV to find the best model with best hyper parameter.
pre-processing method is standard scaler

[52]:

```

lgr_pipeline = Pipeline([("scaler", StandardScaler()), ("LogisticRegression",
    LogisticRegression())])
rfc_pipeline = Pipeline([("scaler", StandardScaler()), ("RandomForestClassifier",
    RandomForestClassifier())])
knn_pipeline = Pipeline([("scaler", StandardScaler()), ("KNeighborsClassifier",
    KNeighborsClassifier())])
gnb_pipeline = Pipeline([("scaler", StandardScaler()), ("GaussianNB",
    GaussianNB())])
sgd_pipeline = Pipeline([("scaler", StandardScaler()), ("SGDClassifier",
    SGDClassifier())])
dt_pipeline = Pipeline([("scaler", StandardScaler()), ("DecisionTreeClassifier",
    DecisionTreeClassifier())])

```

[53]:

```

pip_dict1 ={0:'Logistic Regression' , 1:'RandomForestClassifier' , 2:'KNeighborsClassifier',
    3: 'GaussianNB', 4: 'SGDClassifier', 5: 'DecisionTreeClassifier' }

```

[54]:

```

piplines1=[lgr_pipeline, rfc_pipeline , knn_pipeline, gnb_pipeline , sgd_pipeline ,
    dt_pipeline ]

```

[55]:

```

scores_df = pd.DataFrame(columns = ["Model", "CVScores"])
for i, pipe in enumerate(piplines1):
    score = cross_val_score(pipe, final_features2, target, cv = 10)
    print(pip_dict1[i], ":", score.mean())

```

Logistic Regression : 0.9751217455582278
 RandomForestClassifier : 0.9743850467720389

```
KNeighborsClassifier : 0.9748503332686301
GaussianNB : 0.9795478562150082
SGDClassifier : 0.9774982135396149
DecisionTreeClassifier : 0.972224585735939
```

```
[56]: grid_params = [
    {"classifier": [RandomForestClassifier()],
     "classifier__n_estimators": [50,100,150,200,250,300],
     "classifier__criterion": ["gini", "entropy"],
     "classifier__max_features": ["auto", "sqrt", "log2"],
     },
]
```

```
[57]: pipeline_new = Pipeline([("scaler", StandardScaler()), ("classifier", RandomForestClassifier())])

random_search = RandomizedSearchCV(estimator = pipeline_new,
                                     param_distributions = grid_params, scoring = 'neg_mean_absolute_error',
                                     n_jobs= -1, cv = 8, verbose = 10, random_state = 42)
```

```
[58]: best_model = random_search.fit(X_train, y_train)
```

Fitting 8 folds for each of 10 candidates, totalling 80 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done   5 tasks      | elapsed:  1.1min
[Parallel(n_jobs=-1)]: Done  10 tasks      | elapsed:  1.3min
[Parallel(n_jobs=-1)]: Done  17 tasks      | elapsed:  1.7min
[Parallel(n_jobs=-1)]: Done  24 tasks      | elapsed:  2.0min
[Parallel(n_jobs=-1)]: Done  33 tasks      | elapsed:  2.5min
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:  3.1min
[Parallel(n_jobs=-1)]: Done  53 tasks      | elapsed:  4.1min
[Parallel(n_jobs=-1)]: Done  64 tasks      | elapsed:  4.3min
[Parallel(n_jobs=-1)]: Done  80 out of  80 | elapsed:  5.7min finished
```

```
[59]: best_model.best_params_
```

```
[59]: {'classifier__n_estimators': 300,
       'classifier__max_features': 'log2',
       'classifier__criterion': 'gini',
       'classifier': RandomForestClassifier(max_features='log2', n_estimators=300)}
```

```
[60]: pipeline_rf = Pipeline([('scaler', StandardScaler()),
```

```

('Random Forest Calssifier',  

 RandomForestClassifier(criterion='entropy', max_features='sqrt',  

 n_estimators=150))]

[61]: model = pipeline_rf.fit(X_train, y_train)

[62]: rf_train_predict = pd.DataFrame({'actual' : y_train,  

                                         'predicted' : model.predict(X_train)})  

rf_train_predict.head()

[62]:      actual predicted  

116587      0        0  

36340       0        0  

175763      0        0  

96918       0        0  

71197       0        0

[63]: rf_test_predict = pd.DataFrame({'actual' : y_test,  

                                         'predicted' : model.predict(X_test)})  

rf_test_predict.head()

[63]:      actual predicted  

80120      0        0  

19670       0        0  

114887      0        0  

120110      0        0  

56658       0        0

[64]: print('Accuracy Score for train dataset : ', metrics.  

     accuracy_score(rf_train_predict.actual, rf_train_predict.predicted))  

print('Accuracy Score for test dataset : ', metrics.  

     accuracy_score(rf_test_predict.actual, rf_test_predict.predicted))

Accuracy Score for train dataset :  1.0  

Accuracy Score for test dataset :  0.9962054066031465

[65]: print('ROC-AUC Score for train dataset : ', metrics.  

     roc_auc_score(rf_train_predict.actual, rf_train_predict.predicted))  

print('ROC-AUC Score for validation dataset : ', metrics.  

     roc_auc_score(rf_test_predict.actual, rf_test_predict.predicted))

ROC-AUC Score for train dataset :  1.0  

ROC-AUC Score for validation dataset :  0.9515587050614198

[66]: conn_cm_test = metrics.confusion_matrix(rf_test_predict.actual, rf_test_predict.  

     predicted, [1,0])  

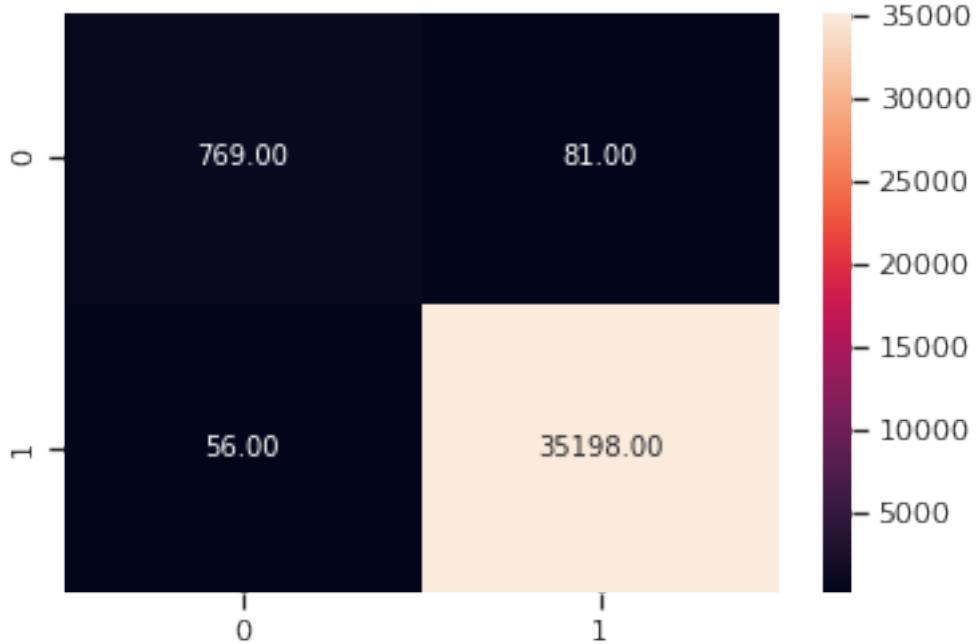
sns.heatmap(conn_cm_test, fmt= '.2f', annot=True)

```

```
/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py:70:  
FutureWarning:
```

```
Pass labels=[1, 0] as keyword args. From version 0.25 passing these as  
positional arguments will result in an error
```

```
[66]: <AxesSubplot:>
```



```
[67]: print(metrics.classification_report(rf_test_predict.actual, rf_test_predict.  
predicted))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	35254
1	0.93	0.90	0.92	850
accuracy			1.00	36104
macro avg	0.96	0.95	0.96	36104
weighted avg	1.00	1.00	1.00	36104

```
[ ]:
```

Forecasting of an order has tardiness

```
[68]: data=df[['Type', 'Benefit per order', 'Sales per customer', 'Delivery Status', 'Customer City', 'Customer Country', 'Customer Segment', 'Customer State', 'Customer Zipcode', 'Department Name', 'Market', 'Order City', 'Order Country', 'Customer Id', 'order date (DateOrders)', 'Order Item Cardprod Id', 'Order Item Discount', 'Order Item Discount Rate', 'Order Item Id', 'Order Item Profit Ratio', 'Sales', 'Order Status', 'Order Item Total', 'Order Profit Per Order', 'Order Region', 'Order State', 'Order Zipcode', 'Product Card Id', 'Product Category Id', 'Product Description', 'Product Image', 'Product Name', 'Product Price', 'Product Status', 'shipping date (DateOrders)', 'Shipping Mode']]
```

```
[69]: features=data.drop(columns=['Late_delivery_risk'])
target=data['Late_delivery_risk']
```

```
[70]: features=features.apply(Labelencoder_feature)
features.head()
```

[70]: Type Benefit per order Sales per customer Delivery Status Category Name Customer City Customer Country Customer Segment Customer State Customer Zipcode Department Name Market Order City Order Country Customer Id order date (DateOrders) Order Item Cardprod Id Order Item Discount Order Item Discount Rate Order Item Id Order Item Profit Ratio Sales Order Status Order Item Total Order Profit Per Order Order Region Order State Order Zipcode Product Card Id Product Category Id Product Description Product Image Product Name Product Price Product Status shipping date (DateOrders) Shipping Mode

0	1		18934		2568		0		40
66			1		0		36		7
4	3		331		70		20649		5961
114			426				4		180516
140	166		2		2568				18934
475			12733		114		47		15
78			78		62		0		0
3									27149
1	3		2272		2559		1		40
66			1		0		36		7
4	3		391		69		19386		1147
114			498				5		179253
44	166		5		2559				2272
841			100315		114		47		13
78			78		62		0		120339
									2209

```

3
2     0          2293          2555          3          40
452      0          0          19385          5          941
4     3          391          69          179252          941
114      529          1          2555          6          1146
44    166          1          114          2293          13
841    100316          62          0          47          120340
78      78          62          0          179252          1980
3
3     1          13638          2546          0          40
285      0          2          19384          5          720
4     3          3226          8          179251          1145
114      600          2          2546          7          13638
119    166          114          47          120341          11
835    100317          62          0          179251          1752
78      78          62          0          13638          11
3
4     2          20599          2526          0          40
66      1          1          19383          36          7
4     3          3226          8          179250          1144
114      682          6          2526          8          20599
156    166          114          47          120342          1528
78      78          62          0          179250          11
3

```

[71]: #Feature Selection

```

# Feature Selection based on importance
from sklearn.feature_selection import f_regression
F_values, p_values = f_regression(features, target)

```

```

/opt/conda/lib/python3.7/site-
packages/sklearn/feature_selection/_univariate_selection.py:302: RuntimeWarning:
invalid value encountered in true_divide

```

[72]: import itertools

```

f_reg_results = [(i, v, z) for i, v, z in itertools.zip_longest(features.
    columns, F_values, [%.3f' % p for p in p_values])]
f_reg_results=pd.DataFrame(f_reg_results, columns=['Variable','F_Value',_
    'P_Value'])

```

[73]: f_reg_results=pd.DataFrame(f_reg_results, columns=['Variable','F_Value',_
 'P_Value'])
f_reg_results = f_reg_results.sort_values(by=['P_Value'])

```
f_reg_results.P_Value= f_reg_results.P_Value.astype(float)
f_reg_results=f_reg_results[f_reg_results.P_Value<0.06]
f_reg_results
```

```
[73]:
```

	Variable	F_Value	P_Value
0	Type	685.999651	0.000
30	Product Description	72.142896	0.000
27	Order Zipcode	25.191994	0.000
36	Shipping Mode	34666.398337	0.000
3	Delivery Status	6798.252477	0.000
25	Order Region	6.848744	0.009
5	Customer City	4.662629	0.031
35	shipping date (DateOrders)	3.556568	0.059

```
[74]: f_reg_list=f_reg_results.Variable.values
f_reg_list
```

```
[74]: array(['Type', 'Product Description', 'Order Zipcode', 'Shipping Mode',
       'Delivery Status', 'Order Region', 'Customer City',
       'shipping date (DateOrders)'], dtype=object)
```

```
[75]: df['Delivery Status'].value_counts()
```

```
[75]:
```

Late delivery	98977
Advance shipping	41592
Shipping on time	32196
Shipping canceled	7754

Name: Delivery Status, dtype: int64

'Delivery Status' should also be omitted. it has directly related to the target

```
[76]: final_features=features[['Type', 'Shipping Mode', 'Order Region',
                           'Customer City', 'shipping date (DateOrders)']]
```

```
[77]: final_data=pd.concat([final_features, target], axis=1)
```

```
[78]: fig = plt.figure(figsize=(20,10))
sns.heatmap(final_data.corr(), annot = True, fmt = '.2f', cmap = 'magma')
```

```
[78]: <AxesSubplot:>
```



```
[79]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(final_features, target, test_size = 0.2, random_state = 42)
```

```
[80]: scores_df = pd.DataFrame(columns = ["Model", "CVScores"])
for i, pipe in enumerate(pipelines1):
    score = cross_val_score(pipe, final_features, target, cv = 10)
    print(pip_dict1[i], ":", score.mean())
```

Logistic Regression : 0.6906918386309326
 RandomForestClassifier : 0.8926758576651675
 KNeighborsClassifier : 0.7423094244520747
 GaussianNB : 0.6718960851534189
 SGDClassifier : 0.6906918386309326
 DecisionTreeClassifier : 0.8855353163083501

```
[81]: grid_params = [
    {"classifier": [RandomForestClassifier()],
     "classifier__n_estimators": [50,100,150,200,250,300],
     "classifier__criterion": ["gini", "entropy"],
     "classifier__max_features": ["auto", "sqrt", "log2"],
     },
    {"classifier": [KNeighborsClassifier()],
     "classifier__n_neighbors": [2,3,4,5,6],
     "classifier__algorithm": ['auto', 'ball_tree', 'kd_tree', 'brute'],
     "classifier__leaf_size": [10,20,30,40,50],
     },
]
```

```
{"classifier": [DecisionTreeClassifier()],
 "classifier__splitter": ["best", "random"],
 "classifier__criterion": ["gini", "entropy"],
 "classifier__max_features": ["auto", "sqrt", "log2"],
 },
]
```

```
[82]: pipeline_new = Pipeline([("scaler", StandardScaler()), ("classifier", RandomForestClassifier())])

random_search = RandomizedSearchCV(estimator = pipeline_new,
                                     param_distributions = grid_params, scoring = 'neg_mean_absolute_error',
                                     n_jobs= -1, cv = 8, verbose = 10, random_state = 42)
```

```
[83]: best_model = random_search.fit(X_train, y_train)
```

Fitting 8 folds for each of 10 candidates, totalling 80 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done   5 tasks      | elapsed:  2.6min
[Parallel(n_jobs=-1)]: Done  10 tasks      | elapsed:  2.6min
[Parallel(n_jobs=-1)]: Done  17 tasks      | elapsed:  2.7min
[Parallel(n_jobs=-1)]: Done  24 tasks      | elapsed:  2.7min
[Parallel(n_jobs=-1)]: Done  33 tasks      | elapsed:  4.3min
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:  4.6min
[Parallel(n_jobs=-1)]: Done  53 tasks      | elapsed:  5.0min
[Parallel(n_jobs=-1)]: Done  64 tasks      | elapsed:  6.7min
[Parallel(n_jobs=-1)]: Done  80 out of  80 | elapsed:  9.2min finished
```

```
[84]: best_model.best_params_
```

```
[84]: {'classifier__n_estimators': 100,
        'classifier__max_features': 'log2',
        'classifier__criterion': 'entropy',
        'classifier': RandomForestClassifier(criterion='entropy', max_features='log2')}
```

```
[85]: pipeline_rfl = Pipeline([('scaler', StandardScaler()),
                               ('RandomForestClassifier', RandomForestClassifier(criterion='entropy'))])
```

```
[86]: model_rfl = pipeline_rfl.fit(X_train, y_train)
```

```
[87]: rfl_train_predict = pd.DataFrame({'actual' : y_train,
                                         'predicted' : model_rfl.predict(X_train)})
rfl_train_predict.head()
```

```
[87]:      actual predicted
116587       1       1
36340        0       0
175763       1       1
96918        1       1
71197        1       1
```

```
[88]: rfl_test_predict = pd.DataFrame({'actual' : y_test,
                                         'predicted' : model_rfl.predict(X_test)})
rfl_test_predict.head()
```

```
[88]:      actual predicted
80120       1       1
19670        1       1
114887       0       1
120110       1       1
56658        0       0
```

```
[89]: print('Accuracy Score for train dataset : ' , metrics.
          accuracy_score(rfl_train_predict.actual, rfl_train_predict.predicted))
print('Accuracy Score for test dataset : ' , metrics.
          accuracy_score(rfl_test_predict.actual, rfl_test_predict.predicted))
```

```
Accuracy Score for train dataset :  0.9999861510230932
Accuracy Score for test dataset :  0.9475958342565921
```

```
[90]: print('ROC-AUC Score for train dataset : ' , metrics.
          roc_auc_score(rfl_train_predict.actual, rfl_train_predict.predicted))
print('ROC-AUC Score for validation dataset : ' , metrics.
          roc_auc_score(rfl_test_predict.actual, rfl_test_predict.predicted))
```

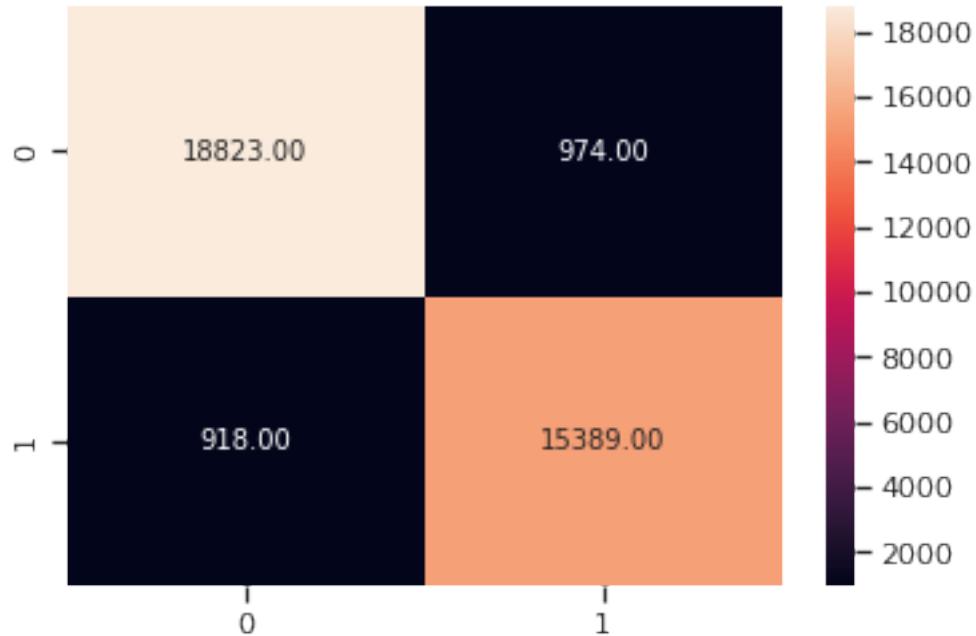
```
ROC-AUC Score for train dataset :  0.9999873705481181
ROC-AUC Score for validation dataset :  0.9472528918259713
```

```
[91]: conn_cm_test = metrics.confusion_matrix(rfl_test_predict.actual, rfl_test_predict.predicted, [1,0])
sns.heatmap(conn_cm_test, fmt= '.2f', annot=True)
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py:70:
FutureWarning:
```

```
Pass labels=[1, 0] as keyword args. From version 0.25 passing these as
positional arguments will result in an error
```

```
[91]: <AxesSubplot:>
```



```
[92]: print(metrics.classification_report(rfl_test_predict.actual,
                                         rfl_test_predict.predicted))
```

	precision	recall	f1-score	support
0	0.94	0.94	0.94	16307
1	0.95	0.95	0.95	19797
accuracy			0.95	36104
macro avg	0.95	0.95	0.95	36104
weighted avg	0.95	0.95	0.95	36104

[]:

11 Forecasting Sales of different products

```
[93]: data_sales=df[['Type', 'Benefit per order', 'Sales per customer',
                    'Delivery Status', 'Late_delivery_risk', 'Category Name', 'CustomerCity',
                    'Customer Country',
                    'Customer Id', 'Customer Segment',
                    'Customer State', 'Customer Zipcode', 'Department Name', 'Latitude',
                    'Longitude',
```

```

    'Market', 'Order City', 'Order Country', 'Order Customer Id', 'order_
    ↴date (DateOrders)', 'Order Id',
    'Order Item Cardprod Id', 'Order Item Discount', 'Order Item Discount_
    ↴Rate', 'Order Item Id',
    'Order Item Product Price', 'Order Item Profit Ratio', 'Order Item_
    ↴Quantity', 'Sales', 'Order Item Total',
    'Order Profit Per Order', 'Order Region', 'Order State', 'Order_
    ↴Status', 'Order Zipcode', 'Product Card Id',
    'Product Category Id', 'Product Description', 'Product Image',_
    ↴'Product Name', 'Product Price', 'Product Status',
    'shipping date (DateOrders)', 'Shipping Mode']]
```

[94]: features=data_sales.drop(columns=['Sales', 'Order Item Quantity', 'Order Item_
 ↴Product Price'])
target=data_sales['Sales']

[95]: features=features.apply(Labelencoder_feature)
features.head()

[95]: Type Benefit per order Sales per customer Delivery Status
Late_delivery_risk Category Name Customer City Customer Country Customer Id
Customer Segment Customer State Customer Zipcode Department Name Latitude
Longitude Market Order City Order Country Order Customer Id order date
(DateOrders) Order Id Order Item Cardprod Id Order Item Discount Order Item
Discount Rate Order Item Id Order Item Profit Ratio Order Item Total Order
Profit Per Order Order Region Order State Order Status Order Zipcode
Product Card Id Product Category Id Product Description Product Image
Product Name Product Price Product Status shipping date (DateOrders)
Shipping Mode

0	1	18934		2568		0		
0		40	66		1	20649		0
36		7		4	3624	4420	3	331
70		20649			5961	65749		114
426			4		180516		140	
2568			18934		15	475	2	
12733		114			47		0	78
78		62	0			27149		3
1	3	2272		2559		1		
1		40	66		1	19386		0
36		7		4	5522	4419	3	391
69		19386			1147	64486		114
498			5		179253		44	
2559			2272		13	841	5	
100315		114			47		120339	78
78		62	0			2209		3
2	0	2293		2555		3		

0	40	452	0	19385	0		
5	941		4	9146	232	3	391
69	19385			1146	64485		114
529		6		179252		44	
2555		2293		13	841	1	
100316		114		47		120340	78
78	62	0				1980	3
3 1		13638		2546		0	
0	40	285		0		19384	2
5	720		4	8467	663	3	3226
8	19384			1145	64484		114
600		7		179251		119	
2546		13638		11	835	2	
100317		114		47		120341	78
78	62	0				1752	3
4 2		20599		2526		0	
0	40	66		1		19383	1
36	7		4	3783	4421	3	3226
8	19383			1144	64483		114
682		8		179250		156	
2526		20599		11	835	6	
100318		114		47		120342	78
78	62	0				1528	3

```
[96]: #Feature Selection based on importance
from sklearn.feature_selection import f_regression
F_values, p_values = f_regression(features, target)

/opt/conda/lib/python3.7/site-
packages/sklearn/feature_selection/_univariate_selection.py:302: RuntimeWarning:
invalid value encountered in true_divide
```

```
[97]: import itertools
f_reg_results = [(i, v, z) for i, v, z in itertools.zip_longest(features.
    ↪columns, F_values, ['%.3f' % p for p in p_values])]
f_reg_results=pd.DataFrame(f_reg_results, columns=['Variable','F_Value', ↪
    'P_Value'])
```

```
[98]: f_reg_results=pd.DataFrame(f_reg_results, columns=['Variable','F_Value', ↪
    'P_Value'])
f_reg_results = f_reg_results.sort_values(by=['P_Value'])
f_reg_results.P_Value= f_reg_results.P_Value.astype(float)
f_reg_results=f_reg_results[f_reg_results.P_Value<0.06]
f_reg_results
```

```
[98]:
```

	Variable	F_Value	P_Value
20	Order Id	1165.171704	0.000
22	Order Item Discount	57166.125441	0.000
21	Order Item Cardprod Id	12782.968321	0.000
39	shipping date (DateOrders)	142.652140	0.000
19	order date (DateOrders)	128.461963	0.000
18	Order Customer Id	673.464036	0.000
27	Order Profit Per Order	13782.670150	0.000
15	Market	240.910781	0.000
28	Order Region	140.517795	0.000
29	Order State	27.935380	0.000
26	Order Item Total	481682.347274	0.000
12	Department Name	524.094617	0.000
32	Product Card Id	12782.968321	0.000
8	Customer Id	673.464036	0.000
33	Product Category Id	10095.780861	0.000
35	Product Image	37751.723070	0.000
5	Category Name	26066.331991	0.000
36	Product Name	37751.723070	0.000
37	Product Price	116680.120560	0.000
2	Sales per customer	481682.347274	0.000
1	Benefit per order	13782.670150	0.000
24	Order Item Id	1133.743612	0.000
16	Order City	8.764240	0.003
9	Customer Segment	4.266892	0.039
31	Order Zipcode	4.200315	0.040

```
[99]: f_reg_list=f_reg_results.Variable.values  
f_reg_list
```

```
[99]: array(['Order Id', 'Order Item Discount', 'Order Item Cardprod Id',  
       'shipping date (DateOrders)', 'order date (DateOrders)',  
       'Order Customer Id', 'Order Profit Per Order', 'Market',  
       'Order Region', 'Order State', 'Order Item Total',  
       'Department Name', 'Product Card Id', 'Customer Id',  
       'Product Category Id', 'Product Image', 'Category Name',  
       'Product Name', 'Product Price', 'Sales per customer',  
       'Benefit per order', 'Order Item Id', 'Order City',  
       'Customer Segment', 'Order Zipcode'], dtype=object)
```

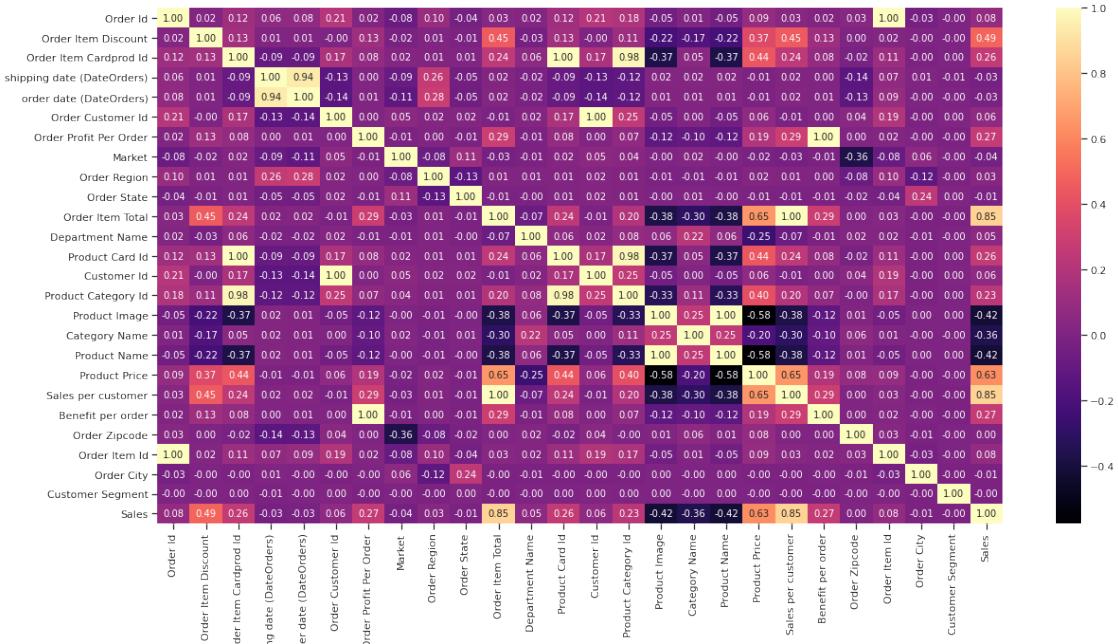
```
[100]: final_features=features[['Order Id', 'Order Item Discount', 'Order Item  
Cardprod Id',  
       'shipping date (DateOrders)', 'order date (DateOrders)',  
       'Order Customer Id', 'Order Profit Per Order', 'Market',  
       'Order Region', 'Order State', 'Order Item Total',  
       'Department Name', 'Product Card Id', 'Customer Id',  
       'Product Category Id', 'Product Image', 'Category Name',
```

```
'Product Name', 'Product Price', 'Sales per customer',
'Benefit per order', 'Order Zipcode', 'Order Item Id',
'Order City', 'Customer Segment']]
```

```
[101]: final_data=pd.concat([final_features, target], axis=1)
```

```
[102]: fig = plt.figure(figsize=(20,10))
sns.heatmap(final_data.corr(), annot = True, fmt = '.2f', cmap = 'magma')
```

```
[102]: <AxesSubplot:>
```



```
[103]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

from sklearn.metrics import *
from sklearn.linear_model import LinearRegression, RANSACRegressor, Lasso, Ridge, SGDRegressor
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
```

```
[104]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(final_features, target,
    test_size = 0.3, random_state = 42)
```

```
[105]: lr_pipeline = Pipeline([("scaler", StandardScaler()), ("linear_regression",
    LinearRegression())])
ridge_pipeline = Pipeline([("scaler", StandardScaler()), ("ridge_regressor",
    Ridge(random_state = 42))])
lasso_pipeline = Pipeline([("scaler", StandardScaler()), ("lasso_regressor",
    Lasso(random_state = 42))])
random_forest_pipeline = Pipeline([("scaler", StandardScaler()), ("randomforest_regression",
    RandomForestRegressor(random_state = 42))])
xgboost_pipeline = Pipeline([("scaler", StandardScaler()), ("xgboost_regression",
    XGBRegressor())])
knn_pipeline = Pipeline([("scaler", StandardScaler()), ("knn_regression",
    KNeighborsRegressor())])
```

```
[106]: pipelines = [lr_pipeline, ridge_pipeline, lasso_pipeline,
    random_forest_pipeline, xgboost_pipeline, knn_pipeline]
```

```
[107]: pipe_dict = {0: "Linear Regression", 1: "Ridge",
    2: "Lasso", 3: "RandomForest", 4: "XGBoost",
    5: "Decision Tree", 6: "KNN"}
```

```
[108]: scores_df = pd.DataFrame(columns = ["Model", "CVScores"])
for i, pipe in enumerate(pipelines):
    score = cross_val_score(pipe, final_features, target, cv = 5)
    print(pipe_dict[i], ":", score.mean())
```

```
Linear Regression : 0.7852647150644007
Ridge : 0.785264849136577
Lasso : 0.7805693158852282
RandomForest : 0.9995414205129233
XGBoost : 0.999651640861491
Decision Tree : 0.9333894519285579
```

```
[109]: grid_params = [
    {"classifier": [XGBRegressor()],
     "classifier_n_estimators": [100,150,200,250,300],
    }
]
```

```
[110]: pipeline_new = Pipeline([('scaler', StandardScaler()), ('classifier',  
    ↪XGBRegressor())])  
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
```

```
[111]: random_search = RandomizedSearchCV(estimator = pipeline_new,  
    ↪param_distributions = grid_params, scoring = 'neg_mean_absolute_error',  
    ↪n_jobs= -1, cv = 8, verbose = 10, random_state = 42)
```

```
[112]: best_model = random_search.fit(X_train, y_train)
```

Fitting 8 folds for each of 5 candidates, totalling 40 fits

/opt/conda/lib/python3.7/site-packages/sklearn/model_selection/_search.py:282:
UserWarning:

The total space of parameters 5 is smaller than n_iter=10. Running 5 iterations.
For exhaustive searches, use GridSearchCV.

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done   5 tasks      | elapsed:  4.7min  
[Parallel(n_jobs=-1)]: Done  10 tasks      | elapsed:  8.3min  
[Parallel(n_jobs=-1)]: Done  17 tasks      | elapsed: 16.5min  
[Parallel(n_jobs=-1)]: Done  24 tasks      | elapsed: 21.5min  
[Parallel(n_jobs=-1)]: Done  33 tasks      | elapsed: 40.3min  
[Parallel(n_jobs=-1)]: Done  38 out of  40 | elapsed: 47.5min remaining:  2.5min  
[Parallel(n_jobs=-1)]: Done  40 out of  40 | elapsed: 47.6min finished
```

```
[113]: best_model.best_params_
```

```
[113]: {'classifier__n_estimators': 300,  
       'classifier': XGBRegressor(base_score=None, booster=None,  
       colsample_bytree=None,  
       colsample_bynode=None, colsample_bytree=None, gamma=None,  
       gpu_id=None, importance_type='gain', interaction_constraints=None,  
       learning_rate=None, max_delta_step=None, max_depth=None,  
       min_child_weight=None, missing=nan, monotone_constraints=None,  
       n_estimators=300, n_jobs=None, num_parallel_tree=None,  
       random_state=None, reg_alpha=None, reg_lambda=None,  
       scale_pos_weight=None, subsample=None, tree_method=None,  
       validate_parameters=None, verbosity=None)}
```

```
[114]: pipeline_XGBRegressor = Pipeline([('scaler', StandardScaler()),  
    ↪('XGBRegressor', XGBRegressor(importance_type='gain', n_estimators=300, ))])
```

```
[115]: model = pipeline_XGBRegressor.fit(X_train, y_train)
```

```
[116]: XGB_train_predict = pd.DataFrame({'actual' : y_train,
                                         'predicted' : model.predict(X_train)})
XGB_train_predict.head()
```

```
[116]:      actual    predicted
99963  149.940002  149.975128
63538  210.850006  210.474564
6661   179.970001  179.993896
93913  129.990005  129.989929
90626  129.990005  129.990906
```

```
[117]: XGB_test_predict = pd.DataFrame({'actual' : y_test,
                                         'predicted' : model.predict(X_test)})
XGB_test_predict.head()
```

```
[117]:      actual    predicted
80120  199.990005  199.987198
19670  250.000000  249.987595
114887 249.899994  249.878433
120110 299.980011  299.969971
56658   119.970001  119.918625
```

```
[118]: predict = model.predict(X_test)
```

```
[119]: r2_score(y_test, predict, multioutput='uniform_average')
```

```
[119]: 0.9974757060247059
```

```
[120]: fig = go.Figure()
fig.add_trace(go.Scatter(x=y_test, y=predict, mode='markers' , name='predicted_vs actual'))
fig.add_trace(go.Scatter(x=y_test , y=y_test, mode='lines' , name='actual'))

fig.update_layout(title='actual Sales vs predicted Sales', xaxis_title= 'Actual Score', yaxis_title = 'Predicted Score' , template= 'plotly_dark')
```

```
[ ]:
```